

Network Monitoring and Management

Introduction to Netflow (Network Flow)

INNOG 6

Dec 8 - 12, 2019




This material is licensed under a Creative Commons Attribution-NonCommercial 4.0 International License (<http://creativecommons.org/licenses/by-nc/4.0/>)

What is a Network Flow

- A set of related packets
- Packets that belong to the same transport connection. e.g.
 - TCP, same src IP, src port, dst IP, dst port
 - UDP, same src IP, src port, dst IP, dst port
 - Some tools consider "bidirectional flows", i.e. A->B and B->A as part of the same flow

[http://en.wikipedia.org/wiki/Traffic_flow_\(computer_networking\)](http://en.wikipedia.org/wiki/Traffic_flow_(computer_networking))

Simple flows



 = Packet belonging to flow X

 = Packet belonging to flow Y

Cisco IOS Definition of a Flow

- Unidirectional sequence of packets sharing:
 - Source IP address
 - Destination IP address
 - Source port for UDP or TCP, 0 for other protocols
 - Destination port for UDP or TCP, type and code for ICMP, or 0 for other protocols
 - IP protocol
 - Ingress interface (SNMP ifIndex)
 - IP Type of Service

which of these six packets are in the same flows?

	<i>Src IP</i>	<i>Dst IP</i>	<i>Protocol</i>	<i>Src Port</i>	<i>Dst Port</i>
A	1.2.3.4	5.6.7.8	6 (TCP)	4001	22
B	5.6.7.8	1.2.3.4	6 (TCP)	22	4001
C	1.2.3.4	5.6.7.8	6 (TCP)	4002	80
D	1.2.3.4	5.6.7.8	6 (TCP)	4001	80
E	1.2.3.4	8.8.8.8	17 (UDP)	65432	53
F	8.8.8.8	1.2.3.4	17 (UDP)	53	65432

which of these six packets are in the same flows? (contd.)

	<i>Src IP</i>	<i>Dst IP</i>	<i>Protocol</i>	<i>Src Port</i>	<i>Dst Port</i>
A	1.2.3.4	5.6.7.8	6 (TCP)	4001	22
B	5.6.7.8	1.2.3.4	6 (TCP)	22	4001
C	1.2.3.4	5.6.7.8	6 (TCP)	4002	80
D	1.2.3.4	5.6.7.8	6 (TCP)	4001	80
E	1.2.3.4	8.8.8.8	17 (UDP)	65432	53
F	8.8.8.8	1.2.3.4	17 (UDP)	53	65432

What about packets “C” and “D”?

Uses and Applications

- You can answer questions like:
 - Which user / department has been uploading / downloading the most?
 - Which are the most commonly-used protocols on my network?
 - Which devices are sending the most SMTP traffic, and to where?
- Identification of anomalies and attacks
- More fine-grained visualisation (graphing) than can be done at the interface level



Uses for Netflow

- Problem identification / solving
 - Traffic classification
 - DoS Traceback
- Traffic Analysis and Engineering
 - Inter-AS traffic analysis
 - Reporting on application proxies
- Accounting (or billing)
 - Cross verification from other sources
 - Can cross-check with SNMP data

Flow Accounting

- A summary of all the packets seen in a flow (so far):
 - Flow identification: protocol, src/dst IP/port...
 - Packet count
 - Byte count
 - Start and end times
 - Maybe additional info, e.g. AS numbers, netmasks
- Records traffic **volume** and **type** but not **content**
- Flow based accounting can be a good supplement to SNMP based accounting.

Flow Accounting (contd.)



Working with flows

1. Configure device (e.g. router) to generate flow accounting records
2. Export the flows from the device (router) to a collector (PC)
 - Configure protocol version and destination
3. Receive the flows, write them to disk
4. Analyse the flows

Many tools available, both free and commercial

Where to generate flow records


1. On a router or other network device

- If the device supports it
- No additional hardware required
- Might have some impact on performance

2. Passive collector (usually a Unix host)

- Receives a copy of every packet and generates flows
- Requires a mirror port
- Resource intensive

Flow Collection

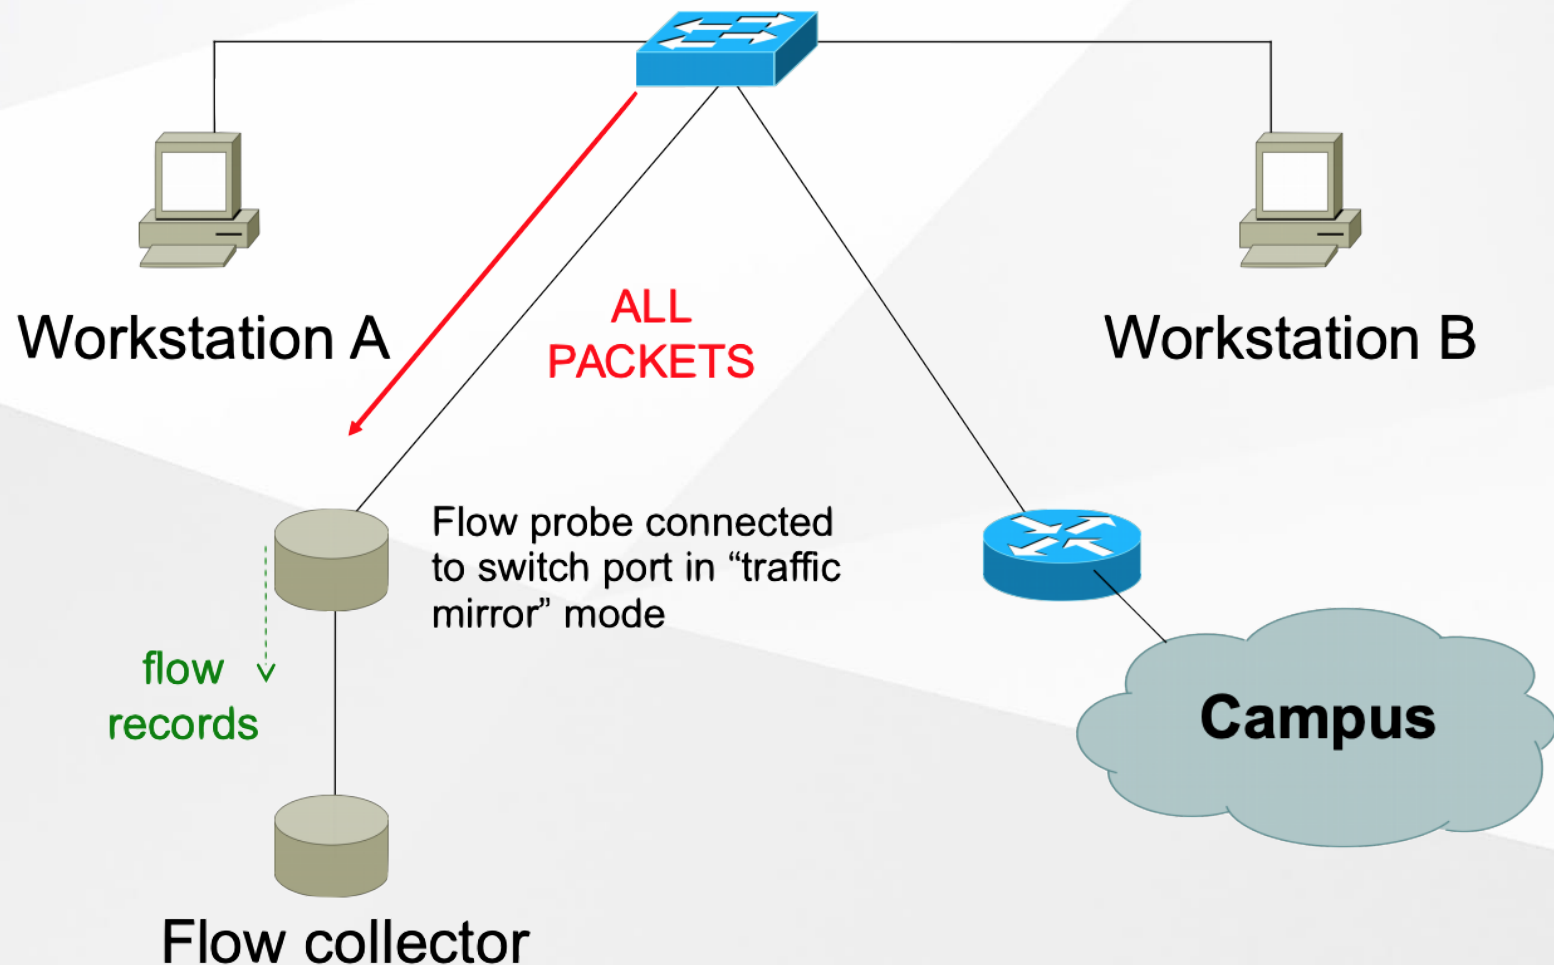


Flow collector
stores exported flows from router.

Flow Collection (contd.)

- All flows through router can be observed
- Router overhead to process & export flows
- Can select which interfaces Netflow collection is needed on and not activate it on others
- If router on each LAN, Netflow can be activated on them to reduce load on core router

Passive Monitor Collection



Passive Collector

- Examples
 - softflowd (Linux/BSD)
 - pfflowd (BSD)
 - ng_netflow (BSD)
- Collector sees all traffic through the network point it is connected on and generates flows
- Relieves router from processing traffic, creating flows and exporting them
- Useful on links:
 - with only one entry into the network
 - where only flows from one section of the network are needed
- Can be deployed in conjunction with an IDS

A thought:

Your network probably already has a device which is keeping track of IP addresses and port numbers of traffic flowing through it.

What is it?

Flow Export Protocols

- Cisco **Netflow**, different versions
 - **v5**: widely deployed
 - **v9**: newer, extensible, includes IPv6 support
- IP Flow Information Export (**IPFIX**):
 - IETF standard, based on Netflow v9
- **sFlow**: Sampling-based, commonly found on switches
- **jFlow**: Juniper
- We use Netflow, but many tools support multiple protocols

Cisco Netflow

- Unidirectional flows
- IPv4 unicast and multicast
 - (IPv6 in Netflow v9)
- Flows exported via UDP
 - Choose a port. No particular standard, although 2055 and 9996 are commonly used
- Supported on IOS, ASA and CatOS platforms
 - but with different implementations

Cisco IOS Configuration

- Configured on each interface
 - Inbound and outbound
 - Older IOS only allows input
- Define the version
- Define the IP address and port of the collector (where to send the flows)
- Optionally enable aggregation tables
- Optionally configure flow timeout and main (v5) flow table size
- Optionally configure sample rate

Configuring Netflow: the old way

- Enable CEF

```
ip cef  
ipv6 cef
```

- Enable flow on each interface

```
ip route cache flow (pre IOS 12.4)
```

or

```
ip flow ingress (IOS 12.4 onwards)  
ip flow egress
```

- Exporting Flows to a collector

```
ip flow-export version [5|9] [origin-as|peer-as]  
ip flow-export destination <x.x.x.x> <udp-port>
```

"Flexible Netflow": the new way

- Only way to monitor IPv6 flows on modern IOS
- Start using it now - IPv6 is here
- Many mind-boggling options available, but basic configuration is straightforward

Flexible Netflow Configuration

- Define one or more exporters

```
flow exporter EXPORTER-1
  destination 192.0.2.99
  transport udp 9996
  source Loopback0
  template data timeout 300
```

- Define one or more flow monitors

```
flow monitor FLOW-MONITOR-V4
  exporter EXPORTER-1
  cache timeout active 300
  record netflow ipv4 original-input
```

```
flow monitor FLOW-MONITOR-V6
  exporter EXPORTER-1
  cache timeout active 300
  record netflow ipv6 original-input
```

Flexible Netflow Configuration (contd.)

- Apply flow monitors to active interface

```
interface GigabitEthernet0/0/0
  ip flow monitor FLOW-MONITOR-V4 input
  ip flow monitor FLOW-MONITOR-V4 output
  ipv6 flow monitor FLOW-MONITOR-V6 input
  ipv6 flow monitor FLOW-MONITOR-V6 output
```


Top Talkers

- You can summarize flows directly on the router, e.g.

```
show flow monitor FLOW-MONITOR-V4 cache aggregate ipv4 source  
address ipv4 destination address sort counter bytes top 20
```

- Yes, that's one long command!
- Old command not available for Flexible Netflow

```
show ip flow top-talkers
```


- Make an Alias:

```
conf t  
alias exec top-talkers show flow..
```

Collecting flows: nfdump

- Free and open source - runs on collector
- nfcapd listens for incoming flow records and writes them to disk (flat files)
 - typically starts a new file every 5 minutes
- nfdump reads the files and turns them into human-readable output
- nfdump has command-line options to filter and aggregate the flows

NfDump Architecture



Date flow start	Duration	Proto	Src IP Addr:Port		Dst IP Addr:Port	Packets	Bytes	Flows
2013-04-18 13:35:23.353	1482.000	UDP	10.10.0.119:55555	->	190.83.150.177:54597	8683	445259	1
2013-04-18 13:35:23.353	1482.000	UDP	190.83.150.177:54597	->	10.10.0.119:55555	8012	11.1 M	1
2013-04-18 13:48:21.353	704.000	TCP	196.38.180.96:6112	->	10.10.0.119:62099	83	20326	1
2013-04-18 13:48:21.353	704.000	TCP	10.10.0.119:62099	->	196.38.180.96:6112	105	5085	1

Analysing flows: nfsen

- Companion to nfdump
- Web GUI
- Creates RRD graphs of traffic totals
- Lets you zoom in to a time of interest and do nfdump analysis
- Manages nfcapd instances for you
 - Can run multiple nfcapd instances for listening to flows from multiple routers
- Plugins available like port tracker, surfmap



Cisco Netflow Versions

Netflow v1

- Key fields: Source/Destination IP, Source/Destination Port, IP Protocol, ToS, Input interface.
- Accounting: Packets, Octets, Start/End time, Output interface
- Other: Bitwise OR of TCP flags.
- Does not have sequence numbers – no way to detect lost flows
- Obsolete

Netflow v2 to v4

- Cisco internal
- Were never released

Netflow v5

- Key fields: Source/Destination IP, Source/Destination Port, IP Protocol, ToS, Input interface.
- Accounting: Packets, Octets, Start/End time, Output interface.
- Other: Bitwise OR of TCP flags, Source/Destination AS and IP Mask.
- Packet format adds sequence numbers for detecting lost exports.
- IPv4 only

Netflow v6 & v7

- Used exclusively on the Cisco Catalyst line of ethernet switches
- Requires the Netflow Feature Card, a specialist forwarding engine for the Catalyst Switches
- Not compatible or comparable with Netflow on Cisco routers

Netflow v8

- Aggregated v5 flows.
- Not all flow types available on all equipment
- Much less data to post process, but loses fine granularity of v5 – no IP addresses.

Netflow v9

- IPv6 support
- 32-bit ASN support
- Additional fields like MPLS labels
- Builds on earlier versions
- Periodically sends "template" packet, all flow data fields reference the template